

Coding LLM Tier List

Every major model ranked by capability, context window, and cost. Built from real harness benchmarks, not marketing pages.

CONTENTS

- 01 About This List
- 02 S-Tier: Best of the Best
- 03 A-Tier: Excellent for Most Tasks
- 04 B-Tier: Solid and Production-Ready
- 05 C-Tier and D-Tier
- 06 Value Picks and Routing Principles

About This List

I personally use only a handful of LLMs day to day, but I have tested around 50 to 60% of these models using the pi.dev harness. This is an ongoing benchmarking effort on my end to evaluate coding LLMs on a common, reproducible harness. The list will iteratively improve as I test more models.

Ranked on three axes: raw coding capability on real tasks, context window size (critical for repo-level and agentic work), and cost per million tokens. Coding-specific fine-tunes consistently outperform same-size general models.

50+

MODELS TESTED

5

TIERS

pi.dev

HARNESS

Jun 2026

LAST UPDATED

S	Best of the Best	Agents, multi-file edits, complex debugging. Use when correctness is non-negotiable.
A	Excellent for Most Tasks	Near-frontier quality at lower cost. Daily coding, production workloads.
B	Solid, Production-Ready	Everyday coding, autocomplete, unit tests. Best fast and cheap options live here.
C	Decent, Lightweight	Simple scripts, boilerplate generation. Struggles with multi-file reasoning.
D	Budget / Edge Deployment	Classification, templating, latency-critical edge use. Not for complex logic.

TIER 02 OF 06

S

Best of the Best

Frontier-class reasoning with 1M+ context windows. These models handle agentic coding loops, complex multi-file edits, and hard debugging problems. Use when correctness matters more than cost.

MODEL	PROVIDER	CONTEXT	INPUT \$/M	OUTPUT \$/M
Claude Opus 4.8 Anthropic	anthropic	1M	\$5.00	\$25.00
Claude Sonnet 4.6 SWEET SPOT Anthropic	anthropic	1M	\$3.00	\$15.00
GPT-5.5 OpenAI	openai	1M	\$5.00	\$30.00
GPT-5.4 OpenAI	openai	1M	\$2.50	\$15.00
Kimi K2.7 Code VALUE MoonshotAI	moonshotai	262K	\$0.95	\$4.00
Qwen3 Coder 480B A35B VALUE Alibaba / Qwen	openrouter	1M	\$0.22	\$1.80
DeepSeek V4 Pro VALUE DeepSeek	deepseek	1M	\$0.44	\$0.87
Grok 4.3 xAI	xai	1M	\$1.25	\$2.50

NOTE DeepSeek V4 Pro and Qwen3 Coder 480B are the steals of this tier. Both sit at S-Tier capability for a fraction of the price of Claude or GPT-5. If you are running agentic loops at scale, start there before defaulting to the obvious names.

TIER 03 OF 06

A

Excellent for Most Tasks

Near-frontier quality at meaningfully lower cost. Strong context handling across the board. Devstral 2 and Qwen3 Coder 30B are standouts: coding-tuned models that punch far above their price point for production workloads.

MODEL	PROVIDER	CONTEXT	INPUT \$/M	OUTPUT \$/M
Claude Opus 4.6 / 4.7 Anthropic	anthropic	1M	\$5.00	\$25.00
GPT-5.2 Codex OpenAI	openai	400K	\$1.75	\$14.00
GPT-5.1 Codex OpenAI	openai	400K	\$1.25	\$10.00
Devstral 2 CODING-TUNED Mistral	mistral	262K	\$0.40	\$2.00
Kimi K2.6 MoonshotAI	moonshotai	262K	\$0.95	\$4.00
Qwen3 Coder 30B A3B VALUE Alibaba / Qwen	qwen	262K	\$0.15	\$0.60
DeepSeek V4 Flash VALUE DeepSeek	deepseek	1M	\$0.14	\$0.28
Gemini 3.1 Pro Preview Google	google	1M	\$2.00	\$12.00
Gemini 3.5 Flash Google	google	1M	\$1.50	\$9.00

NOTE Qwen3 Coder 30B at \$0.15 / \$0.60 is one of the best deals in the entire landscape. Coding-tuned at 262K context for less than most D-Tier general models. DeepSeek V4 Flash (\$0.14 / \$0.28 on a 1M context) is another absurd value at this tier.

TIER 04 OF 06

Solid, Production-Ready

B

Reliable for everyday coding tasks: autocomplete, unit test generation, PR summaries, and code review. Codestral is purpose-built for code completion and still a top choice for that specific job. Claude Haiku 4.5 is the best fast and cheap option in the whole list.

MODEL	PROVIDER	CONTEXT	INPUT \$/M	OUTPUT \$/M
GPT-4.1 OpenAI	openai	1M	\$2.00	\$8.00
Codestral (latest) Mistral	mistral	256K	\$0.30	\$0.90
Devstral Small 2 Mistral	mistral	256K	\$0.10	\$0.30
Qwen3 235B A22B Alibaba / Qwen	qwen	262K	\$0.22	\$0.88
Qwen3 Coder Next Alibaba / Qwen	qwen	262K	\$0.22	\$1.80
Mistral Large 3 Mistral	mistral	262K	\$0.50	\$1.50
Llama 4 Maverick 17B Meta	meta	1M	\$0.24	\$0.97
Claude Haiku 4.5 Anthropic	anthropic	200K	\$1.00	\$5.00
Kimi K2.5 MoonshotAI	moonshotai	262K	\$0.60	\$3.00
MiniMax M3 MiniMax	minimax	512K	\$0.30	\$1.20

NOTE **Devstral Small 2 at \$0.10 / \$0.30 is the best autocomplete model per dollar in the list.** Coding-tuned, 256K context, almost free. Route autocomplete and inline suggestions here before reaching for anything higher.

C Decent, Use for Lightweight Tasks

Good for simple scripts, boilerplate generation, and quick one-off answers. Struggle with complex multi-file reasoning, long dependency chains, or architecture decisions.

MODEL	PROVIDER	CONTEXT	INPUT \$/M	OUTPUT \$/M
GPT-4o OpenAI	openai	128K	\$2.50	\$10.00
Gemini 2.5 Flash Google	google	1M	\$0.30	\$2.50
Gemini 3 Flash Google	google	1M	\$0.50	\$3.00
Qwen3 32B Alibaba / Qwen	qwen	131K	\$0.08	\$0.28
Llama 3.3 70B Meta	meta	131K	\$0.10	\$0.32
Nemotron 3 Super 120B NVIDIA	nvidia	1M	\$0.09	\$0.45
DeepSeek R1 DeepSeek	deepseek	163K	\$0.70	\$2.50
Mistral Medium 3.1 Mistral	mistral	262K	\$0.40	\$2.00

D Budget / Specialized Use Only

Fast and cheap. Best for token classification, simple templating, or latency-critical edge inference. Do not use for multi-step reasoning or anything requiring held context.

MODEL	PROVIDER	CONTEXT	INPUT \$/M	OUTPUT \$/M
GPT-4.1 mini OpenAI	openai	1M	\$0.40	\$1.60
GPT-4.1 nano OpenAI	openai	1M	\$0.10	\$0.40
Mistral Small 3.2 Mistral	mistral	128K	\$0.10	\$0.30
Llama 3.1 8B Meta	meta	131K	\$0.02	\$0.03

06

Value Picks and Routing Principles

BEST OVERALL

DeepSeek V4 Pro

\$0.44 input / \$0.87 output per M tokens

S-Tier capability for a fraction of the flagship price. The default choice if you are cost-sensitive but refuse to drop quality.

BEST CHEAP CODER

Qwen3 Coder 30B A3B

\$0.15 input / \$0.60 output per M tokens

Coding-tuned at 262K context. A-Tier quality at what most D-Tier models charge. Hard to beat for constrained budgets.

BEST FAST AUTOCOMPLETE

Devstral Small 2

\$0.10 input / \$0.30 output per M tokens

Coding-tuned, 256K context, effectively free at scale. Purpose-built for inline completion and fast turnaround tasks.

BEST FREE TIER

Qwen3 Coder 480B (free)

\$0 via OpenRouter free tier

S-Tier model available for free through OpenRouter. Rate-limited but usable for experimentation and low-volume workloads.

BEST LARGE CONTEXT BUDGET

Kimi K2.7 Code

\$0.95 input / \$4.00 output, 262K context

Coding-specific model at 262K context without flagship pricing. Strong for repo-level tasks where context depth matters.

BEST FAST AND CHEAP

Claude Haiku 4.5

\$1.00 input / \$5.00 output per M tokens

Best quality-per-latency tradeoff in the whole list. Route quick tasks here: PR summaries, docstring gen, simple debugging.

ROUTING PRINCIPLES

01 Context window is non-negotiable for agentic work

Prefer 256K+ for large codebases. Models under 128K will fail at real repo-level tasks regardless of how well they score on benchmarks.

02 Coding-tuned models outperform same-size general ones

Always prefer a coding-specific fine-tune when one exists at your price point. Devstral, Codestral, Kimi Code, and Qwen Coder all validate this consistently.

03 Route by task type, not by prestige

S-Tier for complex reasoning and agents. B-Tier for autocomplete. D-Tier for classification and templating. Paying S-Tier rates for boilerplate is just waste.

Prices vary by provider and region. Cross-reference OpenRouter or Amazon Bedrock for routing options. Bedrock adds a ~10% regional markup on S-Tier models. This list updates as more models are benchmarked via the pi.dev harness.